

Name of the Team

Iron Faries

Team Members

Gigi Ma, gigi.ma@zurich.com

Description and quantification of the problem

No one knows how to approach AI — even though everyone wants to automate the 80% repetitive tasks.

- **#StartingHurdle** — 95% lack tech skills; traditional AI needs devs, months, \$\$\$\$. Nothing starts. All they have is unorganized dumps of SOPs, JDs, org charts, workflow diagrams, job shadow recordings.
- **#ExpensiveAI** — "AI = expensive" kills every scaling conversation. No one teaches efficient architecture or how agents think.
- **#OneSizeFailsAll** — Generic chatbots fit no one; staff try once, quit.
- **#TrustGap** — AI repeats mistakes; trust erodes; adoption dies.
- **#ShelfwareRisk** — Role changes = tool obsolete. Rebuild = back to zero.

Cost: \$31K/year inefficient paid hours wasted per person (80% of work hours on manual tasks AI could handle) × 1,000 HK staff = \$31M annual opportunity lost.

What the results of the prototype are

Prototype Results — 1 Built Twin: Sarah Chan, Claims Handler

Each twin will have different metrics; below is our first proof point

Layer 1: Factory Agent (runs once)

- Agents built: 1 twin (Sarah Chan — Claims Handler, 4 skills)
- Token use: ~\$3 one-time
- Output: Twin Agent with 4 skills — Claim Handling, Email Reply, Reporting, EXCO Prep

Layer 2: Twin Agent Daily Run (2-week field trial)

BYO Use Case

Digital Twin Factory — An Agent That Builds Agents

Main learning from experience

- **#ArchitectureFirst** — 3 options evaluated; the "obvious" approach was 6x more expensive. Design structure before features.
- **#AgenticThinking** — Reason → Plan → Decide → Self-Correct. A learnable, repeatable framework — not magic.
- **#GoodAgentCriteria** — Reusability, token efficiency, security, human approval, self-improvement. Without scoring these, agents can't scale.
- **#GuardrailsNotGatekeeping** — Citizen developers CAN touch production data — with HITL, PII masking, audit trails, and parallel runs.
- **#ProgressiveAutonomy** — Start 100% human approval. As accuracy proves, gradually release. Visible risk + improving > invisible risk + no trail.

How your approach is solving the problem

Digital Twin Factory — one solution, five problems killed.

- **#StartingHurdle** → **Agent Builds Agents** — Just talk. Dump SOPs, JDs, org charts, recordings — any format. 11 agents build your twin in 10 min. Zero coding.
- **#ExpensiveAI** → **Cost-Smart 2-Layer** — Claude thinks, n8n does (85% cheaper than Claude alone). Costs less than a coffee a day. Factory teaches efficient architecture and how agents think → costs drop further.
- **#OneSizeFailsAll** → **Your Twin ≠ My Twin** — Same factory, completely different skills per role. Claims twin ≠ Marketing twin ≠ Finance twin — tailored uniquely to you.
- **#TrustGap** → **Easy Feedback, Skills Learn** — Approve / Edit / Reject each output. Every correction permanent. Accuracy rate rises. Edit rate drops.
- **#ShelfwareRisk** → **Grows With You** — New role? Process has changed? The Factory helps you easily add a skill. No rebuild ever.

What are the next steps for scaling your solution

Roll-Out Plan — Champions first, then depth, then breadth

Stage	Goal / Success Looks Like	Who / What
Stage 1 Month 1–2	Prove it works for ANY role on Day 1	2 champions × 8 depts = 16 diverse twins
Stage 2 Month 3–6	Full adoption; champions become peer trainers	ALL users per team. 200+ twins

Problem Approach & Solution Design

One solution, five problems killed — each Winning Theme answers a real pain point

THE PROBLEM


No one knows how to approach AI — even though everyone wants to automate the 80% repetitive tasks.

- **#StartingHurdle** — 95% lack tech skills; needs devs, months, \$\$\$\$. Nothing starts.
- **#ExpensiveAI** — "AI = expensive" kills scaling. No one teaches efficient architecture.
- **#OneSizeFailsAll** — Generic chatbots fit no one; staff try once, quit.
- **#TrustGap** — AI repeats mistakes; trust erodes; adoption dies.
- **#ShelfwareRisk** — Role changes = tool obsolete. Rebuild = zero.

Cost: \$31K/year inefficient paid hours per person × 1,000 HK staff = \$31M lost

THE SOLUTION — DIGITAL TWIN FACTORY

- **#StartingHurdle** → **Agent Builds Agents** — Dump SOPs, JDs, recordings. 11 agents build twin in 10 min. Zero coding.
- **#ExpensiveAI** → **Cost-Smart 2-Layer** — Claude thinks, n8n does (85% cheaper). Teaches efficient architecture.
- **#OneSizeFailsAll** → **Your Twin ≠ My Twin** — Same factory, different skills per role — tailored uniquely to you.
- **#TrustGap** → **Easy Feedback** — Approve / Edit / Reject. Every correction permanent.
- **#ShelfwareRisk** → **Grows With You** — Process changed? Add a skill. No rebuild.



Agent That Builds Agents



5 Winning Themes



Truly Agentic (Level 5)

Results & Key Insights

Prototype: 1 Built Twin — Sarah Chan, Claims Handler (2-week field trial)

Layer 1: Factory Agent (runs once)

1
Twin Built

~\$3
One-time Cost

Output: Twin Agent — Sarah Chan (Claims Handler)
4 Skills: Claim Handling | Email Reply | Reporting | EXCO Prep

Layer 2: Twin Agent Daily Run

Dimension #1: Cost Saving

SKILL	FREQUENCY	COST/DAY	TIME SAVED/WK
Claim Handling	8x/day	\$0.18	5.5 hrs
Email Reply	12x/day	\$0.14	4.5 hrs
Reporting	2x/month	\$0.03	2.5 hrs
EXCO Prep	1x/week	\$0.02	2.0 hrs
TOTAL		\$0.42/day	14.5 hrs/wk

Value: \$31,320/yr | ROI: 204x | Payback: <1 day
Insight: Each twin has unique cost/saving profile depending on role complexity & frequency

Dimension #2: Accuracy

SKILL	ACCURACY
Claim Handling	89%
Email Reply	94%
Reporting	100%
EXCO Prep	88%
OVERALL	91%

Day 1: 84% → Day 7: 91% → Day 14: 95% | Edit rate: 40% → 12%
Insight: HITL makes the greatest change — every correction permanently updates Twin Profile



Cost-Smart Architecture (Dim #1)



Easy Feedback — Skills Learn (Dim #2)

⚠️ 1 built twin (Sarah Chan, Claims). Every twin is different — Marketing = different skills, cost, time-saved. **That's not a limitation; that's #YourTwin≠MyTwin in action.**

Next Steps & Recommendations

Champions first, then depth, then breadth — reusable agents accelerate every stage

ROLL-OUT PLAN

Most diverse twins first → All users per team → Reuse across BUs/regions → Global scale

STAGE 1 — MONTH 1–2

Pilot

Goal: Prove works for ANY role on Day 1

Who: 2 champions × 8 depts = 16 diverse twins

STAGE 2 — MONTH 3–6

Expand HK

Goal: Full adoption, peer training culture

Who: ALL users per team. 200+ twins

STAGE 3 — MONTH 6–12

APAC

Goal: No rebuild from 0, just adapt. Speed ×10

Who: Reuse HK agents across BUs & regions

STAGE 4 — YEAR 2+

Global

Goal: Every employee has a twin

Who: 55,000 staff. \$170M value unlock



Your Twin ≠ My Twin — Same Factory, ∞ different outputs per role

SCALES AFFORDABLY

#CheapToBuild

Low one-time cost. No dev needed. Same Factory, ∞ twins.

#CheapToRun

Minimal daily cost. Built-in coaching → costs drop over time.

#MaxReusability

Built-in agent checks reusability. Max carry-over across roles & regions.



Easy Feedback — Skills Learn & Grow (costs decrease as twin improves)

RISK MITIGATION

Risk 1: AI acts without consent

→ Human-in-the-loop; nothing moves without approval

Risk 2: Data exposure

→ Security-check agent + PII masking

Risk 3: Black box

→ Full audit trail: every decision logged

Risk 4: Disrupts live ops

→ 2-week parallel shadow run before go-live